Unified Training–Learning–Inference–Reinforcement Dynamics Across Neural Systems

Anonymous CVPR submission

Abstract

Nine experiments in nanoGPT first established the open-loop dynamics of Training–Learning–Inference (TLI). Training evolves a parameter–optimizer state: the same realized update can produce different finite-amplitude functional responses in different receiving states. Learning persistently reorganizes distributed functional support, and a target's capability changes when its readout boundary is crossed. Inference recruits and combines learned support at a frozen persistent state.

The response-factor analysis in TL-E05 directly led to a predictor of immediate target-boundary outcomes following realized updates, achieving 91.43% accuracy and 91.49% four-way macro recall on 5,088 nanoGPT confirmation targets. We subsequently tested these TLI relations and the prediction method in five additional neural systems: ResNet-18, diffusion U-Net, routed Vision-MoE, a graph neural network, and a hard-spiking recurrent network.

Finally, in a seventh system, a shared GRU policy with four previously mastered skills, we investigated how concentrating correct positive feedback reorganizes functional support and weakens other mastered capabilities. Redistributing feedback restored their mean accuracy from 70.31% to 99.48% while preserving the reinforced skill.

By linking the consequences of inference to subsequent training actions, Reinforcement closes the experimentally established TLI open-loop dynamics into a unified Training–Learning–Inference–Reinforcement (TLIR) closed-loop dynamics.

1. Introduction

Training, generalization, capability formation, and inference are central problems in neural-network research. Different fields have developed influential theories to study them.

In *optimization theory*, the Neural Tangent Kernel (NTK) describes parameter-induced output changes through the network Jacobian and the corresponding kernel in specified training regimes [8]. Edge of Stability studies training behavior in relation to the largest Hessian eigenvalue, learning rate, and loss trajectory [3]. Sharpness-Aware Minimization (SAM) constructs training updates using loss changes in a parameter neighborhood to improve generalization [6]. Here, the *current state* is represented by parameters and their local derivatives or curvature; *realized updates* are produced by the optimizer; and *responses* appear as changes in model outputs or loss.

In *representation learning and generalization*, Spectral Bias examines frequency-dependent learning rates of functions or kernel modes [17]. Implicit Bias studies why an optimizer selects particular solutions, such as maximum-margin directions on separable data [20]. Neural Collapse measures class-feature means, within-class variability, and classifier geometry, revealing continuing representational change after training accuracy saturates [15]. Information Bottleneck theory studies mutual information among inputs, hidden representations, and targets, including the retention and compression of task-relevant information [21]. Their *internal organizations* take the form of spectral modes, optimizer-selected solutions, feature geometry, or information structure; their *functional changes* appear in learning rates, evolving solutions, and predictive performance.

In *capability formation*, Grokking captures the long delay between fitting training examples and achieving strong test generalization [16]. Subsequent circuit studies track the formation and reorganization of internal computations during this delay [13]. Double Descent studies how test error can rise near interpolation and fall again as model size, training time, or data conditions change [1, 12]. These studies measure the *current state* through training phase or interpolation regime, *internal change* through circuit and solution evolution, and *observable capability* through accuracy and generalization error.

In *continual learning*, catastrophic forgetting occurs when training on a new task reduces performance on earlier tasks. Elastic Weight Consolidation (EWC) estimates parameter importance using Fisher information from earlier tasks and penalizes changes to important parameters [10]. Gradient Episodic Memory (GEM) constrains new-task gradients using their inner products with previous-task gradients [11]. Orthogonal Gradient Descent (OGD)

projects new gradients into directions that approximately preserve earlier-task outputs [5]. In these methods, the *existing state* includes previously learned parameters and function-sensitive directions; *updates* are constrained new-task actions; *interactions* occur through overlapping gradient or functional subspaces; and *capability changes* are measured as retention, forgetting, or transfer.

In *inference research*, Attention uses queries, keys, and values to compute input-dependent information weights [22]. In-context learning studies how examples and task information provided in the context affect predictions while model parameters remain fixed [2]. Chain-of-Thought prompting studies how intermediate reasoning steps affect performance on complex tasks [23]. The *current state* is the trained model; *input conditions* include the query, context, and intermediate computations; *internal computation* varies with attention and activations; and *functional performance* is observed in the resulting predictions.

In *reinforcement and feedback-driven training*, reinforcement learning from human feedback (RLHF) uses preference and reward signals to adjust model behavior [14]. Self-training creates further training data from a model's predictions [24]. Recursive model-collapse studies examine the loss of low-probability modes when successive models are trained on data generated by their predecessors [19]. Here, *feedback* is supplied through rewards, preferences, or generated-data distributions; *training actions* are the resulting updates; and subsequent *functional change* is observed through behavior, accuracy, and generated distributions.

These studies use different mathematical objects—kernels, Hessians, gradients, class features, mutual information, task scores, attention weights, and reward signals—yet six recurring categories can be identified: *current state, realized update, functional response, internal organization, observable capability, and feedback. The core TLI relations were established experimentally before this cross-field comparison; identifying these recurring variables here is a way to help readers understand the theory, not the process by which it was discovered.*

2. Experimental Establishment of TLI Open-Loop Dynamics

The nine nanoGPT experiments that established TLI involved a substantially broader set of native interventions, controlled comparisons, and independent validations than can be presented within the conference page limit. This section summarizes their key experimental findings and the resulting theoretical relations. Complete protocols, execution records, negative results, and independent validators are provided in the accompanying repository.

2.1. Training: Realized Updates and Receiving States

At a registered training step, the parameter–optimizer state is $Z_t = (\theta_t, \omega_t)$, and the realized parameter update is $\Delta\theta_t = \theta_{t+1} - \theta_t$.

TL-E01 revealed capability formation, decline, and recovery that could not be fully characterized by loss, accuracy, or training step alone. TL-E02 further showed that pausing effective parameter–optimizer evolution delays capability formation, while the clipping-threshold control did not reproduce this effect.

TL-E03 used full/skip branches and reciprocal update transplants to establish that the same realized update produces different functional responses in different receiving states:

$$R_q(\theta_A, \Delta\theta_B) \quad \text{vs.} \quad R_q(\theta_B, \Delta\theta_B), \tag{1}$$

where $R_q(\theta, \Delta) = m_q(\theta + \Delta) - m_q(\theta)$. Separate optimizer-memory exchanges examined how historical optimizer states affect update formation:

$$\Delta^{(a,b)} = \mathrm{OptStep}(\theta_t, g_t, \omega_{a,b}). \tag{2}$$

TL-E04 measured the complete finite-amplitude response to an actual update:

$$r_{t,q}(\alpha) = m_q(\theta_t + \alpha\Delta\theta_t) - m_q(\theta_t). \tag{3}$$

The experiments identified near-linear, saturating, accelerating, turnback, and sign-reversal responses. TL-E05 subsequently identified three primary response-conditioning families: target-boundary state (F1), target-specific update geometry (F3), and parameter–optimizer receiving state (F5).

2.2. Learning: Distributed Support and Capability Boundaries

TL-E06 used reversible component interventions to establish distributed functional support, including component necessity, backup, substitution, synergy, and failure slack.

TL-E07 then examined how realized training updates persistently reorganize this support:

$$\Delta P_{t,q} = P_q(\theta_t + \Delta\theta_t) - P_q(\theta_t). \tag{4}$$

Across 72 experimental sections, 1,636 valid support-allocation transitions and 228 primary-support hand-offs were recorded.

TL-E08 linked these internal changes to observable capability through identity-aligned target readout boundaries. The signs of the pre- and post-update margins distinguish maintenance, loss, non-acquisition, and acquisition or recovery. Learning is thus identified with persistent functional reorganization, while capability transitions are determined by crossings of target-specific readout boundaries.

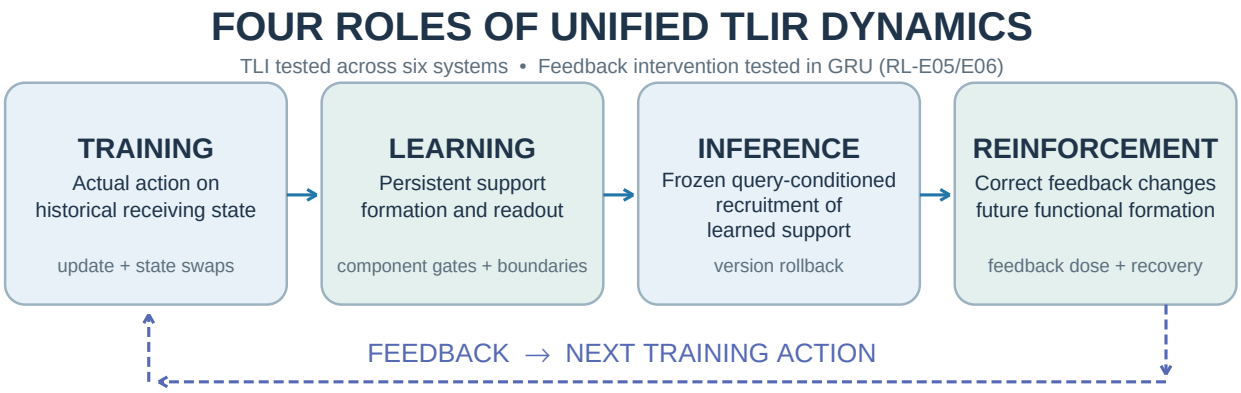


Figure 1. Four peer roles of TLIR. The TLI core is tested in six neural systems; RL-E05/E06 test reinforcement through controlled feedback redistribution in a shared GRU policy.

2.3. Inference: Frozen Recruitment of Learned Support

INF-E01 examined 13 nanoGPT training histories and 52 checkpoints. With persistent learned states frozen, the experiment measured query-conditioned component effects, non-additive support interactions, and changes produced by replacing trained components with their exact pre-formation versions.

At a fixed learned state L^*, inference recruits and combines learned support:

$$S_{q,c}^{\text{active}} = \Pi_{q,c}(L^*), \qquad y_{q,c} = H_{q,c}(S_{q,c}^{\text{active}}). \tag{5}$$

Non-additive component interactions are measured by

$$I_{ij,q} = f_{\text{full}} - f_{-i} - f_{-j} + f_{-i,-j}, \tag{6}$$

while learned-version dependence is tested through

$$D_{j,q} = f(q;\theta^*) - f(q;\theta^{(j\leftarrow\text{pre})}). \tag{7}$$

Exact restoration verified that these interventions did not persistently alter the learned state. The experiments establish Inference as the query-conditioned recruitment and combination of support formed through Learning.

Together, these nine experiments established the connected Training–Learning–Inference relations, forming the TLI open-loop dynamics.

3. Prediction and Cross-System Validation

3.1. Theory-Derived Target-Boundary Prediction

The response-factor analysis in TL-E05 identified three primary conditioning factors: the current target-boundary state (F1), target-specific update geometry (F3), and parameter–optimizer receiving state (F5). These findings led directly to TL-P01, a predictor of immediate target-boundary outcomes following an actual formed update.

For target y, competitor c, and hidden representation h, let $b_c = w_y - w_c$ and $\Delta b_c = \Delta w_y - \Delta w_c$. Using the first and second directional derivatives $h' = Dh[\Delta\theta_t]$ and $h'' = D^2h[\Delta\theta_t, \Delta\theta_t]$, the predictor computes

$$\begin{aligned}\widehat{g}_{y,c} = b_c^\top h + \Delta b_c^\top h + b_c^\top h' \\ + \Delta b_c^\top h' + \tfrac{1}{2} b_c^\top h''.\end{aligned} \tag{8}$$

The target is predicted correct if $\min_{c\neq y} \widehat{g}_{y,c} > 0$. Combining predicted post-update correctness with the known pre-update state gives four capability transitions. Predictions use the already formed update and pre-response information, without reading actual post-update target outputs.

3.2. Cross-System Validation

We subsequently tested the established TLI relations and their predictive consequences in five additional neural systems. Table 1 summarizes the seven systems used in the study, including the separate GRU reinforcement experiment examined in Section 4.

Each of the five additional TLI systems was examined through realized-update interventions, receiving-state comparisons, finite-amplitude responses, component-coalition support measurements, target-boundary outcomes, and frozen-inference tests. Cross-system implementations evaluated all 16 coalitions of four registered components, with exact Shapley effects:

$$\phi_{j,q} = \sum_{C\subseteq\mathcal{C}\setminus\{j\}} \frac{|C|!(3-|C|)!}{4!} \left[v_q(C\cup\{j\}) - v_q(C)\right]. \tag{9}$$

Here $v_q(C)$ denotes the target margin under active component coalition C.

System (optimizer)	Task / registered components
nanoGPT (Adam)	Cyclic next-token; 4 attention/MLP branches
ResNet-18 (SGD momentum)	CIFAR-100 classification; 4 residual stages
Diffusion U-Net (AdamW)	CIFAR-10 denoising; 4 U-Net paths
Vision-MoE (AdamW)	CIFAR-100 classification; 4 routed experts
GNN (RMSprop)	Graph reachability; 4 message-passing blocks
Spiking SNN (Adagrad)	8×8 digits; 4 spiking blocks
GRU (RL-E05/E06; AdamW)	Four-skill feedback; 4 shared hidden groups

Table 1. Seven neural systems: six TLI chains and a GRU feedback study.

All 42 ResNet, 18 diffusion U-Net, and 36 Vision-MoE parameter receiving-state exchanges exceeded their registered response-difference thresholds. Finite-amplitude nonlinear responses occurred in 32.23%, 19.25%, and 35.71% of their respective target-update records, while support reallocation exceeded its specified threshold in 91.41%, 97.02%, and 80.43%. Their frozen-inference experiments further confirmed query-conditioned support, non-additive interactions, and dependence on exact trained component versions.

The GNN reproduced the main TLI relations across three seeds, although one seed failed the registered nonlinear-response frequency criterion (1.56%). The hard-spiking SNN passed all twelve registered TLI criteria across four seeds, extending the evidence to discontinuous spiking dynamics.

The prediction algorithm was also evaluated across these systems, with task-specific adaptations of the directional response computation. Table 2 reports immediate post-update target-boundary classification.

The predictor achieved high four-way recall in five systems. In the hard-spiking SNN, discontinuous activation thresholds limited local derivative-based prediction, yielding 70.74% macro recall despite 95.31% accuracy. Together, the experiments demonstrate cross-system recurrence of the core TLI relations and identify architecture-dependent differences in their quantitative responses and predictive approximations.

System	Accuracy (%)	Four-way recall (%)
nanoGPT	91.43	91.49
ResNet-18	99.28	97.99
Diffusion U-Net	99.60	97.98
Vision-MoE	99.63	96.95
GNN[a]	98.54	98.62
Hard-spiking SNN[a]	95.31	70.74

[a] Independent prediction tests use registered GNN/SNN training histories.

Table 2. Post-update target-boundary classification reports. Macro recall averages four observed transition categories; high accuracy alone is not sufficient. Predictor details and evaluation sets differ by system.

4. Reinforcement: From TLI Open-Loop to TLIR Closed-Loop Dynamics

4.1. Theory-Derived Prediction of Feedback Dynamics

The established TLI relations show that training persistently reorganizes distributed functional support, while inference recruits and combines this learned support to produce behavior. When the consequences of inference are repeatedly fed back into subsequent training, the resulting updates can further reorganize the support on which future inference depends.

This led to a prediction: sustained, narrowly concentrated correct positive feedback could progressively favor the reinforced capability while weakening other already mastered capabilities. The predicted effect arises from feedback-directed functional reorganization, rather than incorrect rewards.

4.2. Experimental Tests and Results

RL-E05/E06 tested this prediction in a seventh system, a shared attention-free GRU policy. All four skills reached mastery before the experimental branches began. Every branch started from an identical parameter–AdamW state, and each complete two-decision episode received $r = 1$ only when correct, and $r = 0$ otherwise.

RL-E05 compared selective, balanced, and frozen feedback. RL-E06 further varied the allocation of 60 episodes per update over 3,200 updates (Table 3).

RL-E05 observed relative support concentration, cross-skill interference, and component-version-dependent tradeoffs. Its stronger preregistered temporal-precedence criterion was not met.

In RL-E06, all 12 formal seeds exhibited dose-ordered changes. The reinforced skill's mean support share increased from 25.54% under balanced feedback to 35.20%

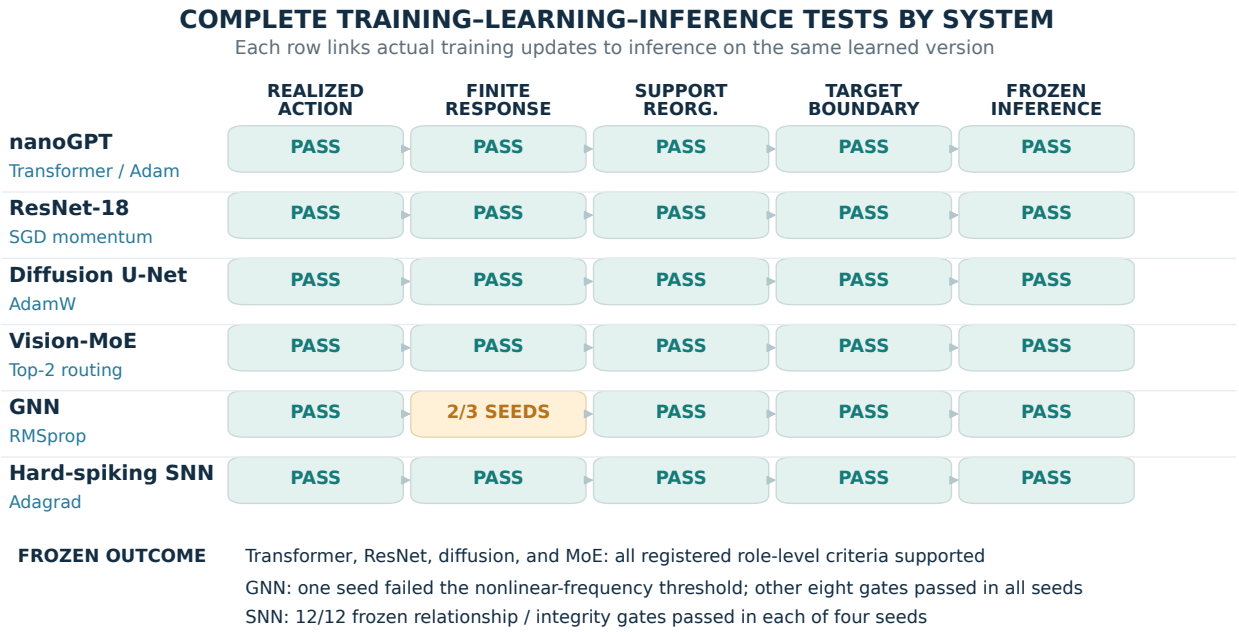


Figure 2. Within-system training–learning–inference experiment chains. One GNN seed meets eight of nine registered gates, with nonlinear-response incidence below its specified threshold; SNN passes all twelve registered gates in four seeds.

Feedback	q_0	q_1	q_2	q_3
Balanced	15	15	15	15
Mild	30	10	10	10
High	45	5	5	5
Exclusive	60	0	0	0

Table 3. RL-E06 feedback episodes per update, by capability.

under exclusive feedback, while the other skills' mean margins decreased from 9.474 to 3.962. Mild and high concentration already reduced these margins without reducing accuracy. Exclusive feedback decreased other-skill accuracy by 39.06 percentage points relative to balanced feedback.

Two recovery branches then started from the exact exclusive-feedback state at update 800. Rebalanced feedback restored other-skill accuracy from 70.31% to 99.48% while preserving the reinforced skill at 100%. Training only the other three skills restored their accuracy to 99.83%, but reduced the previously reinforced skill to 89.06%.

These experiments establish that, in the tested shared system, the distribution of correct feedback can redirect persistent functional support, alter previously mastered capabilities, and reverse these changes through feedback redistribution.

4.3. TLIR Closed-Loop Dynamics

Reinforcement connects the consequences of inference to subsequent training actions. The resulting training updates reorganize learned support, which in turn changes future inference and its consequences. This feedback closure transforms the experimentally established TLI open-loop dynamics into the unified Training–Learning–Inference–Reinforcement (TLIR) closed-loop dynamics.

5. Unified Explanations of Positive and Negative Findings

Question and discovery order. A separate comparison of 21 historical positive/negative pairs asks whether common receiver–action–response–support–boundary–feedback relations explain both the success of local accounts and the findings that limit their scope. The examples below illustrate this shared interpretation.

Kernel dynamics versus feature learning. In the corresponding infinite-width lazy regime, NTK dynamics give a rigorous local training description [8]; feature-learning parameterizations admit different limiting behavior [25]. TLI locates both within receiver- and parameterization-dependent action–response regimes: a valid frozen-kernel projection does not imply that all finite-amplitude training trajectories preserve their features. The two results occupy distinct training regimes.

Sharp minima versus useful sharpness-aware training. Function-preserving reparameterizations can change raw parameter sharpness without changing generalization [4]. Nevertheless, SAM improves performance in specified tasks by optimizing a neighborhood-robust loss [6]. The paired explanation is that raw parameter geometry is not the complete learned functional state, while a sharpness-aware intervention can still change the *actual training action* beneficially. Thus a non-invariant explanatory scalar and an effective local control method can coexist.

Compression as a useful mechanism, not a universal requirement. General compression-stage claims fail in some architectures and measurement regimes [18], whereas appropriately controlling information bottlenecks can yield valid generalization guarantees [9]. Persistent functional reorganization in TLI does not require a monotonic decrease of $I(X;T)$, but compression can be one regime-specific outcome of the receiving state, noise, and objective. This predicts *bounded utility without universal necessity*.

Delayed emergence and continued learning after mastery. In grokking, test capability can appear abruptly after training performance has saturated, although internal computational structure evolves in identifiable stages [13]. Neural-collapse results likewise document continuing feature geometry after training error vanishes [15]. TLI separates persistent support formation from the discrete target readout boundary. Continued internal formation and sudden observable transitions are thus compatible outcomes of the same measured distinction.

Forgetting versus preservation and recovery. Continual learning can damage older capabilities, while methods such as EWC protect selected high-importance parameter directions [10]. In TLI, the signed response of an old target depends on overlap between the new action and its function-sensitive support, relative to remaining boundary slack. The GRU feedback forks supply a matched intervention example: concentrated *correct* feedback erodes other mastered skills, whereas rebalanced feedback largely restores them.

Beneficial self-training versus recursive collapse. Pseudo-label-based self-training improves accuracy in suitable data mixtures [24]; recursively replacing original data with model-generated samples can instead erase minority tails [19], while accumulating real data can prevent that degradation in studied regimes [7]. Both follow a feedback-data distinction: an inference-generated sample is neither intrinsically beneficial nor intrinsically harmful; its retained information and subsequent training action matter. For an under-retained mode with retention factor $0 < \rho < 1$, a conditional illustrative recursion is

$$q_{t+1} = (1-\alpha)\rho q_t + \alpha q_{\text{real}}. \tag{10}$$

At $\alpha = 0$ the mode contracts, whereas $\alpha > 0$ supplies a positive fixed-point floor if $q_{\text{real}} > 0$. The recursion models the case of a repeatedly under-retained data mode.

Theory-first extrapolation. Beyond historical comparisons, the TLIR relations have been used to derive conditional predictions and their test conditions before searching the relevant literature. Examples include post-capability max-margin selection under exponential-tail losses [20] and evolving-state stability at the Edge of Stability [3]. These comparisons identify specific state, action, and boundary conditions under which local mathematical results arise. Future tests can directly intervene on these conditions and evaluate the predicted transitions.

6. Conclusion

Through systematic experiments across neural systems, this work reveals dynamical relations spanning neural-network training, learning, inference, and reinforcement. From capability formation and use to loss and recovery, TLIR establishes experimentally testable causal connections between observable performance changes and the organization and evolution of internal functional support. Neural-network capability evolution thus moves beyond a black-box performance phenomenon and becomes a dynamical process that can be investigated through precise experimental interventions and quantitative prediction.

LLM Usage Statement

OpenAI ChatGPT was used for research ideation and discussion, assistance with experimental-protocol design and code development/review, interpretation of experimental records, and manuscript drafting, editing, and LaTeX formatting. The authors are responsible for the study design, all scientific claims, and the reported results.

References

[1] M. Belkin, D. Hsu, S. Ma, and S. Mandal. Reconciling modern machine-learning practice and the classical bias–variance trade-off. *PNAS*, 116(32):15849–15854, 2019. 1

[2] T. B. Brown et al. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 2020. 2

[3] J. M. Cohen et al. Gradient descent on neural networks typically occurs at the edge of stability. *Proc. ICLR*, 2021. 1, 6

[4] L. Dinh, R. Pascanu, S. Bengio, and Y. Bengio. Sharp minima can generalize for deep nets. *Proc. ICML*, pages 1019–1028, 2017. 6

[5] M. Farajtabar, N. Azizan, A. Mott, and A. Li. Orthogonal gradient descent for continual learning. *Proc. AISTATS*, pages 3762–3773, 2020. 2

[6] P. Foret, A. Kleiner, H. Mobahi, and B. Neyshabur. Sharpness-aware minimization for efficiently improving generalization. *Proc. ICLR*, 2021. 1, 6

[7] M. Gerstgrasser et al. Is model collapse inevitable? breaking the curse of recursion by accumulating real and synthetic data. *ICML Workshop on Foundation Models in the Wild*, 2024. 6

[8] A. Jacot, F. Gabriel, and C. Hongler. Neural tangent kernel: Convergence and generalization in neural networks. *Advances in Neural Information Processing Systems*, 2018. 1, 5

[9] K. Kawaguchi et al. How does information bottleneck help deep learning? *Proc. ICML*, 2023. 6

[10] J. Kirkpatrick et al. Overcoming catastrophic forgetting in neural networks. *PNAS*, 114(13):3521–3526, 2017. 1, 6

[11] D. Lopez-Paz and M. A. Ranzato. Gradient episodic memory for continual learning. *Advances in Neural Information Processing Systems*, pages 6467–6476, 2017. 1

[12] P. Nakkiran et al. Deep double descent: Where bigger models and more data hurt. *Proc. ICLR*, 2020. 1

[13] N. Nanda, L. Chan, T. Lieberum, J. Smith, and J. Steinhardt. Progress measures for grokking via mechanistic interpretability. *International Conference on Learning Representations*, 2023. 1, 6

[14] L. Ouyang et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 2022. 2

[15] V. Papyan, X. Y. Han, and D. L. Donoho. Prevalence of neural collapse during the terminal phase of deep learning training. *PNAS*, 117(40):24652–24663, 2020. 1, 6

[16] A. Power, Y. Burda, H. Edwards, I. Babuschkin, and V. Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets, 2022. arXiv:2201.02177, 2022. 1

[17] N. Rahaman et al. On the spectral bias of neural networks. *Proc. ICML*, pages 5301–5310, 2019. 1

[18] A. M. Saxe et al. On the information bottleneck theory of deep learning. *Proc. ICLR*, 2018. 6

[19] I. Shumailov et al. Ai models collapse when trained on recursively generated data. *Nature*, 631:755–759, 2024. 2, 6

[20] D. Soudry et al. The implicit bias of gradient descent on separable data. *J. Mach. Learn. Res.*, 19(70):1–57, 2018. 1, 6

[21] N. Tishby and N. Zaslavsky. Deep learning and the information bottleneck principle, 2015. in IEEE Information Theory Workshop, 2015, arXiv:1503.02406. 1

[22] A. Vaswani et al. Attention is all you need. *Advances in Neural Information Processing Systems*, pages 5998–6008, 2017. 2

[23] J. Wei et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in Neural Information Processing Systems*, 2022. 2

[24] Q. Xie, M.-T. Luong, E. Hovy, and Q. V. Le. Self-training with noisy student improves imagenet classification. *Proc. CVPR*, pages 10687–10698, 2020. 2, 6

[25] G. Yang and E. J. Hu. Tensor programs iv: Feature learning in infinite-width neural networks. *Proc. ICML*, pages 11727–11737, 2021. 5